

HCLTech AMPLIFIED ROUND 2
SOLUTION DOCUMENTATION

LearnPath

AI-Powered Personalized Learning Path Recommender

"Tell us your goal. We'll find your gaps and build the path - not someone else's."

Team submission - github.com/Jathin-24/LearnPath

1. Problem Understanding & Solution Design

- Online learning platforms recommend individual courses, but learners struggle to identify the right **sequence** of resources to reach a specific goal.
- Different learners have different skill levels, interests, career aspirations and learning preferences - a one-size-fits-all approach is ineffective.
- Learners waste hours covering concepts they already know and miss the prerequisites they actually need.

LearnPath addresses this with a closed-loop adaptive system: conversational intake -> profiling -> deterministic assessment -> personalized sequenced roadmap -> learning -> reassessment -> adaptation. The design explicitly rejects "an LLM generates a roadmap" in favor of **an adaptive learning operating system that continuously decides what to learn next** based on evidence.

Direct mapping to the brief: conversational interface (Profiler Agent), learner profiling engine (structured LearnerProfile), recommendation engine (two-path RAG + LLM selection), personalized path generator with prerequisites & milestones, an AI explainer ("why this recommendation"), and a progress dashboard with next recommended actions.

2. The Two-Path Model (Deep Dive)

The recommendation engine is a **hybrid two-path model** that blends dataset grounding with live web expertise. It is the core intellectual contribution of the project.

Aspect	Path A - Dataset	Path B - Open Web
Source of truth	Enriched 80-course, 109,776-review dataset (static, curated)	Live Tavily web + YouTube search (current, broad)
Candidates	TF-IDF + FAISS nearest-neighbour retrieval (k=15)	Tavily top-4 web + top-3 YouTube results
Selection	LLM picks semantic fit from retrieved candidates	LLM synthesizes notes & plans topics (3-6)
Ordering	Prerequisite graph + topological sort (Kahn's algorithm)	Sequential chain (each node prereqs the prior)
Grounding	Course name, summary, concepts, difficulty, review_count	Real URLs with relevance score, freshness, reason
Env. footprint	Free, offline, no model download (free-tier memory)	Requires TAVILY_API_KEY
Best for	Goals covered by the 80-course catalog	Novel / emerging / niche goals

2.1 Routing & confidence filter

The decision of "which path" is made inside `run_path_a()` in `backend/agents/path_a.py`, governed by three constants:

```
SEED_K = 15 # top-N FAISS candidates retrieved
MIN_SIMILARITY_SCORE = 0.35 # cosine-sim threshold a candidate must pass
TEMPLATE_SIMILARITY_THRESHOLD = 0.85 # cached-roadmap reuse threshold
```

- **Query build** - the learner's goal plus every detected skill gap joined with ". " (fallback: "general programming fundamentals").
- **Template cache check** - the query is embedded and matched against prior learners' roadmaps; at cosine ≥ 0.85 a very similar prior roadmap is reused wholesale (nodes re-locked), skipping all retrieval/LLM work.
- **Retrieve + filter** - 15 candidate courses are retrieved; only those scoring ≥ 0.35 survive.
- **Whole-goal fallback to Path B** - if **zero** candidates pass 0.35, the module logs `dataset_match_too_weak` and routes the entire roadmap to Path B.

The 0.35 threshold is evidence-based: live testing on "become a backend developer" showed real matches scoring 0.36-0.59 while noise (Android App Development, React Native, even Ethical Hacking Basics) sat at 0.28-0.32 - no boundary a human would trust from raw vector math, so the LLM refines selection on top.

Crucially, **both paths can be combined in one roadmap**: when a selected Path-A course lists *external prerequisite concepts* (knowledge outside the catalog, e.g. "HTTP Basics"), those become Path-B stub nodes embedded in the same roadmap and are wired as real prerequisite edges. The result is a dataset-grounded roadmap whose external gaps are filled from the open web.

3. Path A - Dataset-Grounded Agent (Deep Dive)

File: `backend/agents/path_a.py`. Responsibility: *RAG retrieval + LLM planning + prerequisite traversal over the 80-course dataset - dataset-grounded roadmap nodes.*

Step 1 - Retrieve candidates

- The goal+gap query is embedded with the TF-IDF vectorizer (4096 features) and searched against the FAISS inner-product index (L2-normalized $\Rightarrow$ cosine similarity).
- **k = SEED_K = 15** candidate courses are returned, score-sorted; those below the 0.35 confidence floor are dropped.

Step 2 - LLM semantic planning

- Each surviving candidate is sent to the LLM as `{course_name, difficulty, concepts[:8], summary}` together with the learner profile (goal, timeline, interests, optional roadmap_instructions).
- The prompt explicitly instructs: *"Exclude anything off-topic... It's fine to select fewer courses than were given - never pad the list."*
- Output is validated against `RoadmapPlanOutput`; names outside the candidate set are filtered; at least one valid name is required; on failure it retries once with a stricter prompt, then fails loud (never silently wrong).
- **This is the semantic-fitness step** - TF-IDF is lexical, so the LLM supplies the semantic judgement about genuine relevance.

Step 3 - Prerequisite graph traversal

- **Closure** - an iterative stack walk expands each selected course's `internal_prerequisites` into the full prerequisite closure set.
- **Topological sort (Kahn's algorithm)** - in-degree counting with a sorted zero-in-degree queue (alphabetical tie-break) guarantees prerequisites appear before dependents; a cycle raises a `clear ValueError`.

Step 4 - Node construction & external promotion

- Each node carries `node_id` (slugified course name), `topic`, `path_type` = `PATH_A_DATASET`, `course_name`, `course_search_link`, `course_summary`, `internal_prerequisites` (slugified, only if in the selected set), `external_prerequisite_concepts`, and the first 5 concept tags.
- External prerequisite concepts hit by selected courses are **promoted into Path-B stub nodes** placed first in the roadmap and added to each dependent's prerequisite list - turning dead links into live, fillable web nodes.

4. Path B - Open-Web Agent (Deep Dive)

File: `backend/agents/path_b.py`. Responsibility: *Tavily web/YouTube search + LLM synthesis - open-web roadmap nodes, external-prerequisite fills, and "find more resources" refresh.*

```
QUESTIONS_PER_NODE = 3
MAX_WEB_RESULTS = 4
MAX_YOUTUBE_RESULTS = 3
```

Search

- Web: `"{topic} tutorial guide"`, top 4 results.
- YouTube: `"{topic} tutorial"`, top 3 results, domains restricted to `youtube.com /youtu.be`.

Synthesis

- Each result's title/URL + first 800 chars are fed to an LLM (`max_tokens=700`) to produce a 200-400 word plain-prose study guide; degrades gracefully to "use general knowledge, note that clearly" if no results exist.
- Every resource records provenance: `source_type` (web/youtube), a 220-char snippet, Tavily score (3 decimals), published date, and an auto-generated *reason* - e.g. "Top result for this topic from {domain}; published {freshness}; Tavily relevance {score:.0%}."

Three fill shapes

Shape	Trigger	What it does
<code>_fill_resources_only</code>	Brand-new stub, no notes yet	Search + synthesize notes; project/quiz deferred
<code>_supplement_node</code>	Node already has resources ("Find more")	Refresh notes/sources/links only; never touches project/quiz
<code>_fill_node</code>	<code>force=True</code> (explicit Regenerate)	Resources AND project/quiz together

Standalone mode (whole-goal Path B)

- For goals with no dataset coverage, the agent searches `"{goal} roadmap topics to learn in order"`, then an LLM lists 3-6 sequential topics.
- Bare stub nodes are chained (each prereqs the previous), then filled later by the Roadmap Generator's per-node loop calling back into `run_path_b(node_id=...)`.

5. The Dataset & Data Pipeline (Deep Dive)

5.1 The raw dataset

Raw dataset: 109,776 Udemy reviews across 80 courses, with no metadata. Each row is a review tied to a course name (course_reviews.csv: Index, Reviews, Course). The metadata - concepts, difficulty, strengths, weaknesses - had to be *engineered*, not assumed.

5.2 One-time enrichment (enrich_courses.py)

- Reviews are grouped by course; at most **10 reviews per course** are sampled (MAX_REVIEWS_PER_COURSE_SAMPLE).
- **One LLM call per course** (max_tokens=1000) extracts: *concepts* (5-12 specific technical concepts), *difficulty* (beginner|intermediate|advanced), *strengths*, *weaknesses*, and a one-sentence *summary*. Parse failures are recorded rather than crashing.
- A **search_link** is built per course: `udemy.com/courses/search/?q={course_name}`.
- `review_count` = number of raw reviews for that course.
- One additional LLM call across the full list (max_tokens=4000) infers the **prerequisite DAG** (course -> prerequisite courses), guided by concept overlap and difficulty progression; the prompt nudges most courses toward 0-2 prerequisites.
- **Total cost: ~81 LLM calls, run once.**

5.3 Internal vs external split (split_prerequisites.py)

- For each course: `internal_prerequisites` = prerequisites that exist as courses in the catalog; `external_prerequisite_concepts` = those that don't.
- External concepts are exactly what Path A promotes to Path-B stub nodes.

5.4 The validated output (enriched_courses.json)

Verified profile of the enriched dataset:

- **80 courses**; total **109,776 reviews**.
- **106 internal prerequisite edges** and **13 external prerequisite concepts** (across 10 courses).
- Difficulty mix: **56 intermediate / 18 beginner / 6 advanced**.
- **0 parse errors**; 80/80 have summary + search_link; 0 empty concept lists; the assertion test verifies every `internal_prerequisites` reference resolves to a real node earlier in the roadmap (0 dangling edges).

```
"React Native Mobile Development": {
  "concepts": ["JSX components", "StyleSheet", "Flexbox layout", "..."],
  "difficulty": "intermediate",
  "strengths": "Reviewers consistently praise the practical, hands-on assignments...",
  "weaknesses": "They repeatedly note a desire for more complex projects...",
  "summary": "Overall, the course is seen as a solid, value-for-money...",
  "search_link":
    "https://www.udemy.com/courses/search/?q=React+Native+Mobile+Development",
  "review_count": 1402,
  "internal_prerequisites": ["JavaScript Fundamentals", "React.js Development"],
  "external_prerequisite_concepts": []
}
```

6. RAG Index & Retrieval (Deep Dive)

6.1 Index build (backend/rag/build_index.py)

- Each course's embedding text = "{course_name}. Difficulty: {difficulty}. Concepts: {concepts}. {summary}".
- Vectorizer: TfidfVectorizer(max_features=4096, sublinear_tf=True, norm="l2") - **4096 features**.
- Embeddings (dense float32) -> faiss.IndexFlatIP (flat inner-product). Because vectors are L2-normalized, inner product = **cosine similarity**.
- Artifacts committed to the repo: course_index.faiss, course_metadata.json (80 names), tfidf_vectorizer.pkl.

6.2 Query-time retrieval (backend/rag/retriever.py)

- retrieve(query, k) loads cached index/vectorizer/courses, transforms and L2-normalizes the query, runs index.search, and returns the <=k results with enriched fields, descending by cosine score.
- embed_text() returns the 4096-dim vector, reused by Path A for the template-similarity cache.
- A clear FileNotFoundError guides the operator if the index is missing.

6.3 Why the TF-IDF approach?

Render free-tier constraint (the deciding factor): the earlier stack bundled torch + sentence-transformers, which blew past the free tier's build/startup memory budget. TF-IDF has no model download and no heavy ML dependency, so the same 80-course index builds, loads and serves comfortably on Render's 512MB free tier. The accepted tradeoff - TF-IDF is lexical - is compensated by the hybrid planner letting the LLM do semantic candidate selection on top of retrieval. No paid embedding API anywhere.

7. Multi-Agent Orchestration & Roadmap Flow (Deep Dive)

Orchestrated by LangGraph (backend/orchestrator/graph.py). The compiled build chain is **PROFILER -> ASSESSMENT -> PATH_A -> ROADMAP_GENERATOR -> END**, with two routing functions: route_entry trusts next_agent to pick where a resumed session starts, and route_next checks awaiting_input first (so a paused node routes straight to END, preventing infinite self-loops) then cascades to next_agent.

The 8 agents (backend/agents/)

Agent	File	Responsibility
Profiler	profiler.py	Conversational onboarding intake -> updates LearnerProfile (goal, timeline, interests, skills) in the same LLM call as the reply.
Assessment	assessment.py	Two-phase skill-gap check - concept checklist + adaptive MCQ -> SkillGapMap; grading is deterministic, never LLM-as-judge.
Path-A	path_a.py	RAG retrieval + 0.35 confidence filter + LLM planning + prerequisite-graph traversal -> dataset-grounded nodes (falls back to Path B if nothing fits).
Path-B	path_b.py	Tavily web/YouTube search + LLM synthesis -> open-web nodes, external-prerequisite fills, "find more resources" refresh.
Roadmap Generator	roadmap_generator.py	Merges Path-A/B output into the final Roadmap, fills web-node resources eagerly, owns template-cache writes and lazy project/quiz/subtopic generation.
Knowledge Extractor	knowledge_extractor.py	Turns imported context / resume text into structured knowledge-base facts (LLM path with regex fallback for resumes).
Explainer	explainer.py	On-demand "why is this topic in my roadmap?" grounded in profile data + prerequisite-graph position (/roadmap/explain/{node_id}).
Tutor	tutor.py	Mid-roadmap Topic Tutor answering questions grounded in the current node (/chat).

Generate -> review -> confirm

- **Generate** - Roadmap Generator splits nodes into dataset vs web; eagerly fills every PATH_B_OPEN_WEB stub's resources; sets stage = ROADMAP_REVIEW and pauses the graph for user review. It also writes a skeleton to the roadmap_templates cache (goal_text, TF-IDF goal_embedding, nodes) so a similar future goal can reuse it - skeleton only, never project/quiz content.
- **Review** - the learner can reorder (locked nodes), skip, add, modify (re-runs Path-A/B with roadmap_instructions) or delete before committing.
- **Confirm** - POST /roadmap/confirm sets stage = IN_PROGRESS and unlocks **exactly one node at a time** in topological order; Path-B nodes without notes are never unlocked.

Lazy generation (don't spend LLM calls on unreachable modules)

- ensure_subtopics -> 4-8 sub-concepts per node, first AVAILABLE, rest LOCKED.
- generate_subtopic_quiz -> 2-question quiz per sub-concept on "Done Learning".
- generate_final_content -> ONE combined LLM call per node for a ProjectAssignment (title, description, success criteria) + TopicAssessment (3 questions) + estimated_days, fired only when every subtopic is PASSED/SKIPPED. Path A uses _generate_node_content; Path B uses generate_project_and_quiz_from_notes (no re-search).

8. LLM Client & Failover (Deep Dive)

```
LLM_PROVIDERS=groq,groq,openrouter,xai # priority order, comma-separated
LLM_API_KEYS=key1,key2,key3,key4 # position-matched
LLM_MODELS=model1,model2,model3,model4 # one per endpoint
DATABASE_URL=postgresql://localhost/learning_path_db
TAVILY_API_KEY=...
```

- Priority-ordered endpoints; lengths of providers/keys/models must be equal (validated at load).
- Endpoint states: **HEALTHY**, **COOLING_DOWN** (rate-limited, skipped until cooldown), **DEAD** (non-retryable, skipped for the run); cooling endpoints auto-re-promote.
- Error classification: 429/"rate limit"/"quota" -> cooling 60s; 401/403/"invalid api key" -> DEAD; >=500/"timeout"/"connection" -> transient retry same endpoint with $1.0 \cdot 2^{\text{attempt}}$ backoff (max 2), then next endpoint.
- All endpoints exhausted -> AllEndpointsExhaustedError caught by the API, which returns a friendly 502 instead of crashing.
- Providers: xai, groq, openrouter (OpenAI-compatible) + native anthropic. status() exposes a key-redacting /health view (label = provider:model:...KEY_LAST4).

8b. Architecture & Stack

Runtime flow: React/Vite -> FastAPI -> LangGraph -> RAG (FAISS + TF-IDF) -> multi-provider LLM / Tavily -> PostgreSQL.

- **Frontend:** React 19, TypeScript, Vite, Tailwind v4; React Flow (roadmap) + Recharts (analytics).
- **Backend:** Python 3.12, FastAPI, Pydantic v2, LangGraph; eight agents; PostgreSQL JSONB per session (no ORM); tables auto-created.
- **RAG:** scikit-learn TF-IDF + FAISS over the enriched 80-course dataset.
- **LLM:** multi-provider (Groq + OpenRouter + xAI) with automatic failover; Tavily for web search.
- **Data pipeline:** enrich_courses.py (one-time LLM pass) + split_prerequisites.py -> validated enriched_courses.json (0 parse errors).

9. Local Setup & Execution

Prerequisites: Python 3.11+, Node.js 18+, PostgreSQL running locally (or a Supabase URL).

Backend (from the repo root):

- python -m venv .venv and activate it (.venv\Scripts\activate on Windows)
- pip install -r requirements.txt
- copy .env.example .env -> fill LLM_PROVIDERS / LLM_API_KEYS / LLM_MODELS / DATABASE_URL / TAVILY_API_KEY (ACCESS_TOKEN_SECRET optional in dev)
- Create the database: CREATE DATABASE learning_path_db; (tables auto-create on first run)
- (Optional) rebuild RAG index: python -m backend.rag.build_index - only if enriched_courses.json changes
- Run: uvicorn backend.api.main:app --port 8000

Frontend:

- cd frontend
- npm install
- npm run dev - open http://localhost:5173

The frontend works out of the box against http://127.0.0.1:8000. Only if your backend lives elsewhere, set VITE_API_BASE_URL in frontend/.env.

Health/status: http://127.0.0.1:8000/health - API docs: http://127.0.0.1:8000/docs

Windows note: do not use --reload - the WatchFiles reloader orphans the worker process. Restart manually.

10. Testing, Verification & Benchmark

- **Backend:** pytest backend/tests/ -v (60+ tests, real Postgres + real LLM, no mocking).
- **AI benchmark:** python -m backend.benchmarks.ai_benchmark --limit 5 (--resume after rate limits; --offline for deterministic heuristics). The hybrid engine is scored against a generic-LLM baseline and a RAG-only baseline over 30 learner profiles across 5 quality dimensions (relevance, ordering, personalization, feasibility, grounding).
- **Frontend:** npm run lint (oxlint) and npx tsc --noEmit.
- **Reliability:** deterministic grading, append-only progress log, multi-key LLM failover, validated LLM output, 0 dangling prerequisite edges.
- **Security:** bcrypt password hashing; signed access tokens (ACCESS_TOKEN_SECRET) enforce session/resource ownership; Guest flow stays open.

11. Deployment (if available)

Piece	Host	Notes
Database	Supabase	Session pooler URL; tables auto-create on startup
Backend	Render	uvicorn backend.api.main:app --host 0.0.0.0 --port \$PORT
Frontend	Vercel	Root directory: frontend; framework Vite

Full configuration (env vars, CORS via ALLOWED_ORIGINS, wake-up notes) is in docs/deployment_guide.md. RAG artifacts are committed - no build step at deploy.

Appendix - Did we build what was asked?

The HCLTech AMPLIFIED brief asked for a conversational interface, learner profiling, a recommendation engine, a personalized path generator with prerequisites and milestones, an AI explainer, and a progress dashboard. Every one of these is implemented and integrated, and we went beyond with adaptation (roadmap regeneration), per-topic quizzes and projects, resource provenance, engagement features, and an AI quality benchmark.

Brief requirement	Built?	Where
Conversational interface (goal in natural language)	Yes	Profiler Agent intake + Chat
Learner profiling engine	Yes	LearnerProfile, resume/context import
Recommendation engine (courses / projects / resources)	Yes	Two-path RAG + LLM + Tavily

Brief requirement	Built?	Where
Personalized path generator w/ prerequisites & milestones	Yes	Roadmap Generator, topological order
AI assistant - explains each recommendation	Yes	Explainer Agent + Topic Tutor
Dashboard - progress, skills, milestones, next actions	Yes	Dashboard + Analytics
Adapt suggestions based on feedback/progress	Yes	Roadmap regeneration, benchmark